



Autonomous Search With AI

Group 28



Current Speaker:
Patrick Bauer

Group Members & Roles



Patrick
Bauer

Team Lead
Simulation Environment
Navigation



Pablo
Trivino

Object Detection
Robot Sensors



Mark
Pedroso

Object Detection
AI Training



Noah
Avizemer

Simulation Environment
Mapping



Nathanael
Cassagnol

Navigation

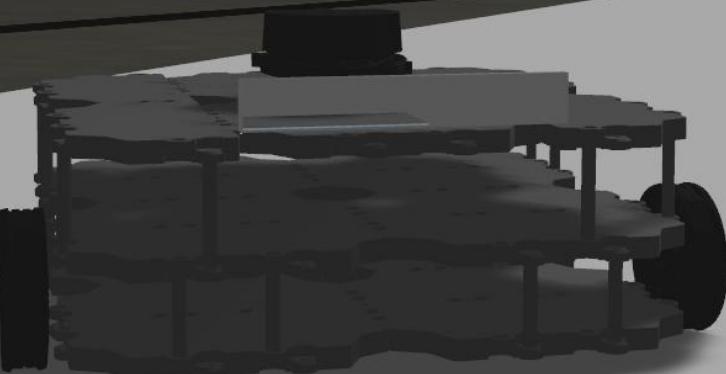
**Sponsor:
Joseph Rivera**



Current Speaker:
Patrick Bauer

Project Objective

Simulate a fully autonomous robot that searches for a target in an urban environment using AI.



Requirements - Simulation

- Simulated urban environment (400ft x 400ft)
 - Several streets with buildings, trees, cars and other obstacles.
- Fully autonomous navigation
 - After the simulation starts, no human assistance.
- GPS denied environment (no global coordinates)
- Must notify user when target is found and simulation is complete.

Requirements - Object Detection

- Bounding box shall accurately surround target
- Must accurately identify the target
 - > 0.7 confidence.
- Model predictions shall be correct at least 97% of the time.
- Model shall make predictions at least 2 times per second.



Current Speaker:
Mark Pedroso



bb8



Implementation

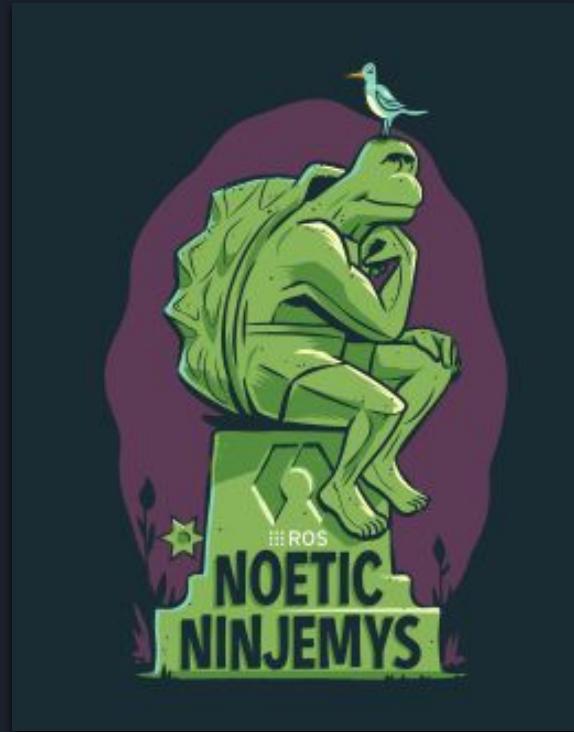
- Software
 - ROS Noetic
 - Gazebo 9
- Algorithms
 - SLAM
 - YOLO



Current Speaker:
Patrick Bauer

ROS

- Open-source robotics operating system
- Consists of infrastructure, tools, and capabilities for programming robots
- Industry standard for programming robot behavior
- Capability to implement packages and algorithms such as YOLO and SLAM



Current Speaker:
Patrick Bauer



Gazebo

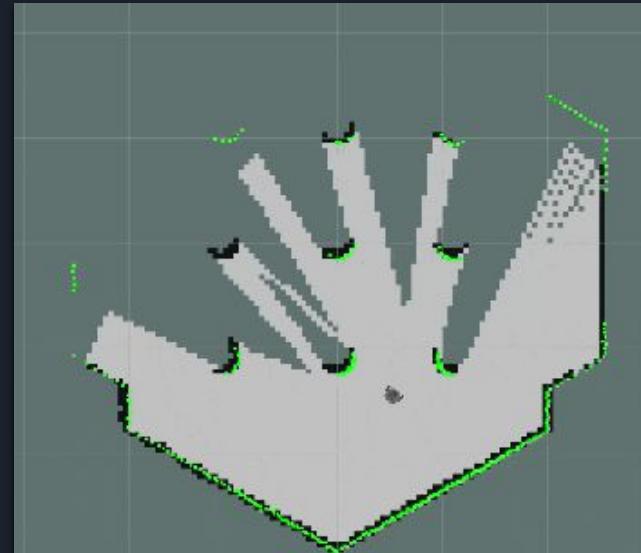
- Open-source robotics simulation software
- Simulate robotic vehicles in various 3D environments
- Supports various sensors and controls for the robot that are necessary for the simulation
- Easy integration with ROS



Current Speaker:
Patrick Bauer

Simultaneous Localization and Mapping (SLAM)

- Allows for an agent to navigate through an environment, avoid obstacles, and generate a map using LIDAR sensors.
- Uses occupancy grid-based mapping



Current Speaker:
Noah Avizemer

You Only Look Once (YOLO)

- Object detection algorithm
- Great precision/fps ratio
- It is a single shot detector
 - Useful for real-time object detection
- Training can be done on standard GPU
- Detects one object per cell
- Can generalize objects well

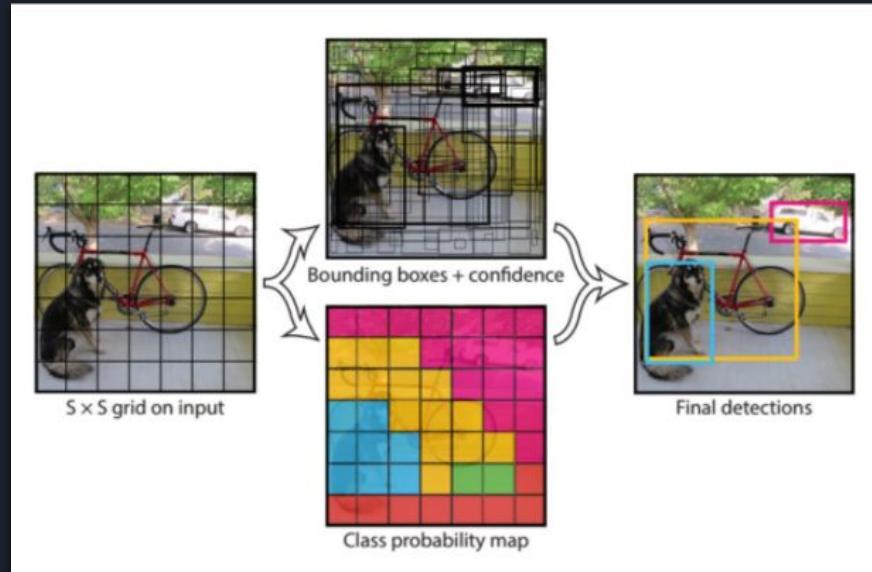


Figure 1. Redmon, J. (2018)



Current Speaker:
Mark Pedroso

Object Detection Consideration

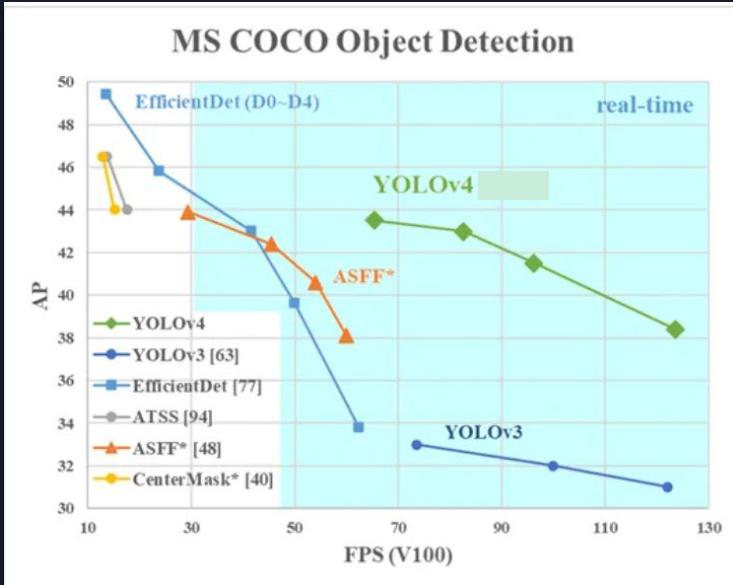


Figure 2. Bochkovskiy, A. (2020)

- YOLOv4 seemed best suited for this project considering it had the best ratio of time to precision
- 10% and 12% increase in AP respectively from YOLOv3

Final Decision for Object Detection

- Yolov3 is more established and stable
 - Yolov3 - 2018
 - Yolov4 - 2020
- Yolov3 has more available documentation for implementation

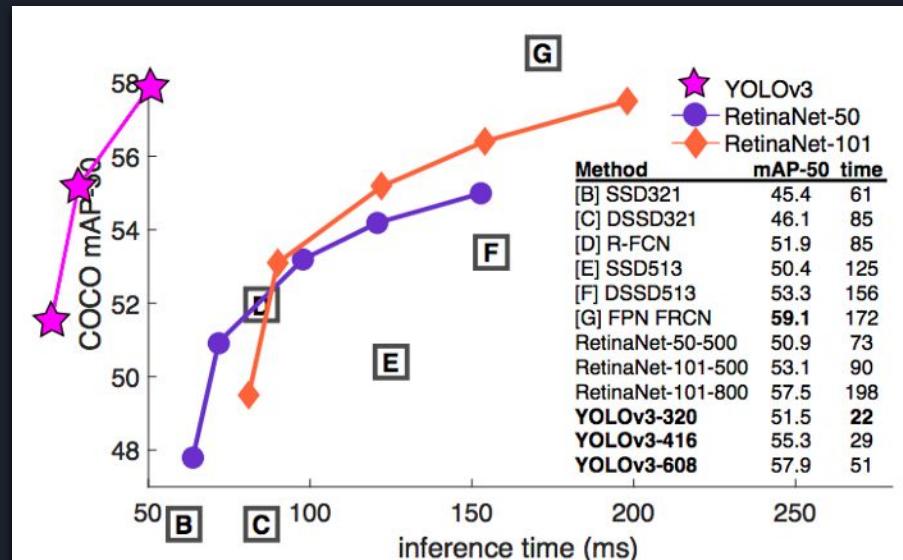


Figure 3. Redmon, J. (2018)

Custom Dataset

YOLOv3 requires a dataset of images to learn

- 1,344 BB8 images gathered for training
 - located within the simulation.
 - Varying backgrounds
 - Different angles & distances
- Labeled using labellImg
 - Created ground truth boxes, and labeled them with the appropriate class
- Dataset split into training and testing sets
 - 80% training, 20% testing



Current Speaker:
Pablo Trivino



Training Methodology

- Model trained in Google Colab using the custom dataset
 - Training time: ~4 hours
- Modified hyperparameters
 - Batches: 64
 - Subdivisions: 16
 - Width/height: 416 x 416
 - Max Batches (2000 * class number): 2000
 - Steps (80%, 90% of max batches): 1600, 1800
- Data augmentation used to improve training
 - Built in image resizing in YOLO config file
 - Dataset contains images from many angles, positions, and backgrounds
- Model tested using the testing set after training
- Final weights file generated by the training was downloaded and used for detections in our local machine



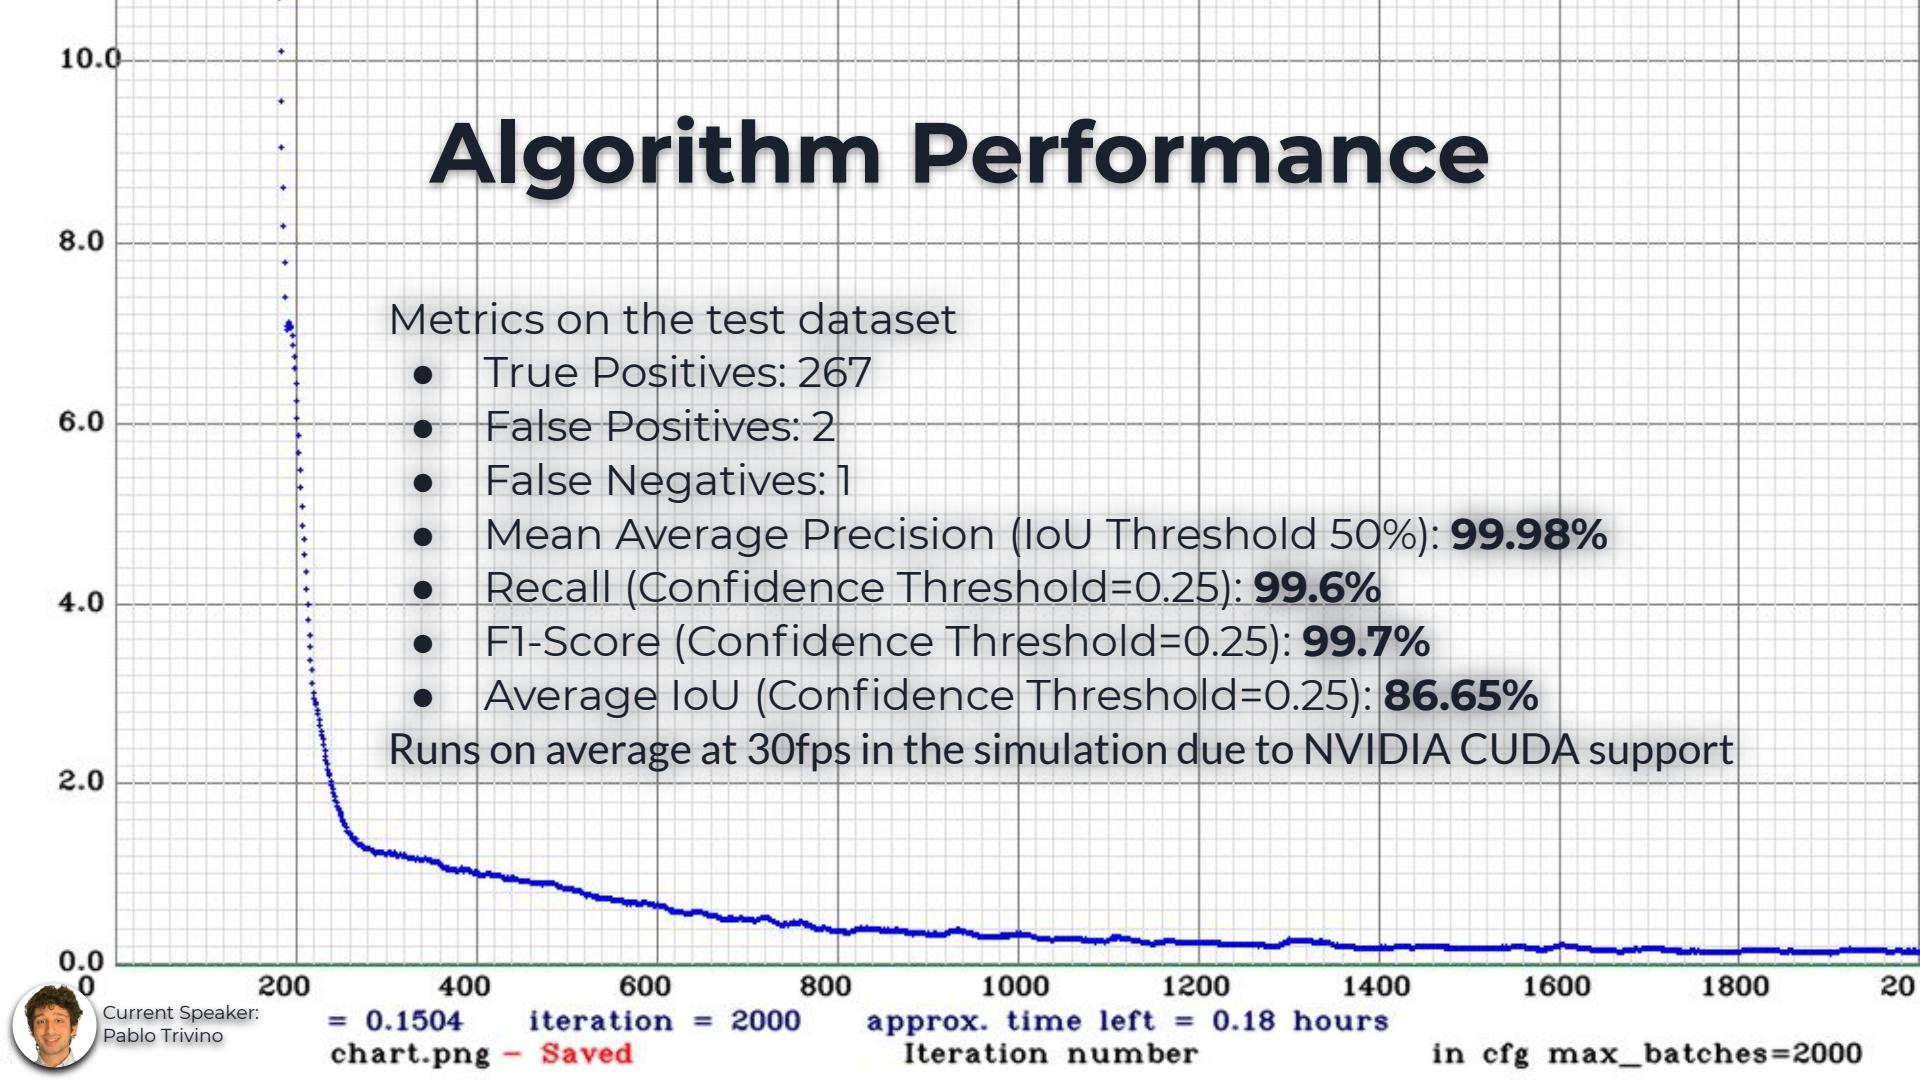
Current Speaker:
Pablo Trivino

Algorithm Performance

Metrics on the test dataset

- True Positives: 267
- False Positives: 2
- False Negatives: 1
- Mean Average Precision (IoU Threshold 50%): **99.98%**
- Recall (Confidence Threshold=0.25): **99.6%**
- F1-Score (Confidence Threshold=0.25): **99.7%**
- Average IoU (Confidence Threshold=0.25): **86.65%**

Runs on average at 30fps in the simulation due to NVIDIA CUDA support



Simulation - Environment

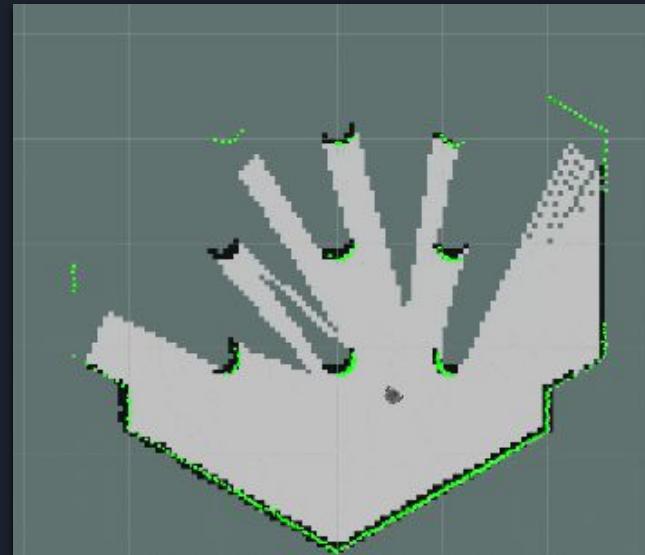
- 400ft x 700ft urban environment
- Lots of different obstacles
- Simulated pavements and grass have different friction attributes to simulate real world conditions.



Current Speaker:
Noah Avizemer

Navigation - Gmapping & SLAM

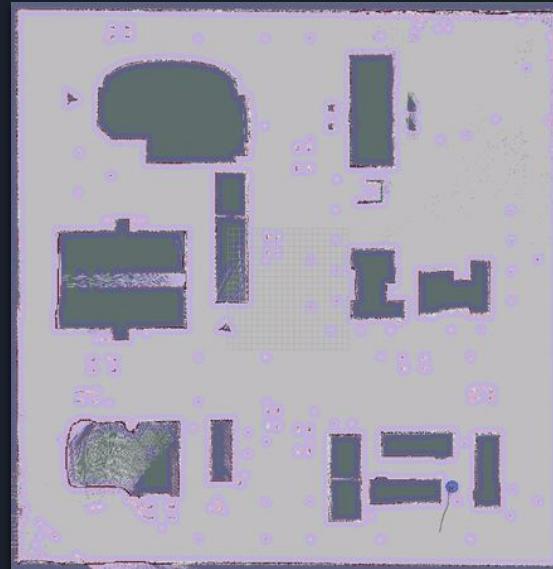
- Map of environment created using Gmapping SLAM
 - Costmap is generated from lidar
- Turtlebot autonomously navigates to areas with unmarked pixels to explore the environment.
- Exploration reveals new areas to turtlebot allowing for the map to grow.



Current Speaker:
Patrick Bauer

Navigation - Grid Search

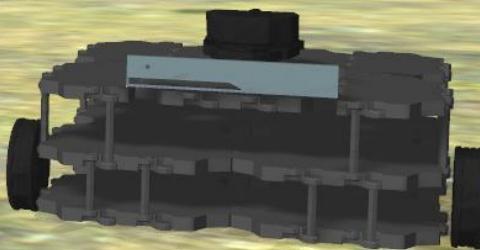
- At runtime waypoints are generated for each point on the grid.
- Modular, can be fit to any environment dimensions.
- A* Algorithm used to navigate to each waypoint.
- If a waypoint cannot be reached, the turtlebot moves to the next one.
- Spaced optimally ensure turtlebot thoroughly searches the entire environment.



Current Speaker:
Nathanael Cassagnol

Navigation - Obstacle Avoidance

- Map-based, A*, pathfinding used to plan paths around obstacles.
- Lidar used to avoid nearby obstacles.
- Avoids getting stuck, damaging itself, or running into anything else.



Current Speaker:
Nathanael Cassagnol

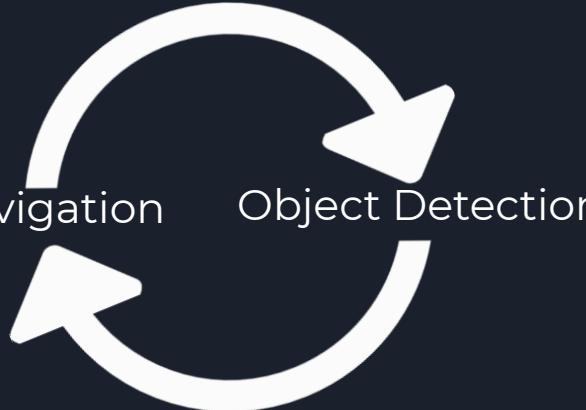
Navigation - Target Approach

- Camera feed is run through our trained AI and the results are fed to navigation.
- If target object detected (>0.6 probability), grid search is paused.
- Turtlebot approaches target based off the target's bounding box.
- If target confirmed mission ends successfully.
- If false positive occurs, grid search continues.



Current Speaker:
Nathanael Cassagnol

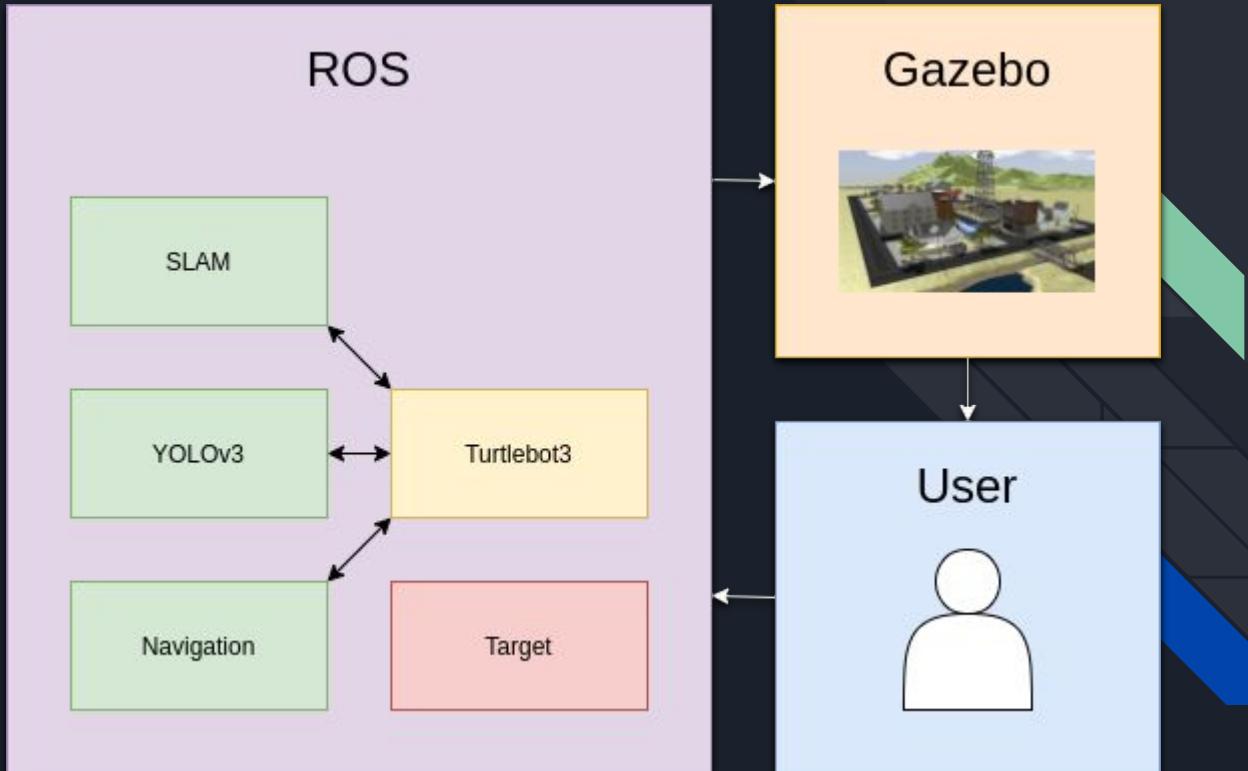
Simultaneous Navigation & Detection

- Need to navigate the environment and search for the object simultaneously.
 - 2 python threads
 - One controls the navigation of the turtlebot.
 - One checks for objects detected.
 - Shared resources for notifying each other about status
- 
- The diagram consists of two white circular arrows on a dark background. The top arrow points clockwise and is labeled 'Navigation'. The bottom arrow points counter-clockwise and is labeled 'Object Detection'. The two arrows are positioned such that they overlap slightly, indicating a continuous loop between the two processes.



Current Speaker:
Patrick Bauer

Entity Relationship Diagram: Project



Current Speaker:
Patrick Bauer

Mission Process Flowchart

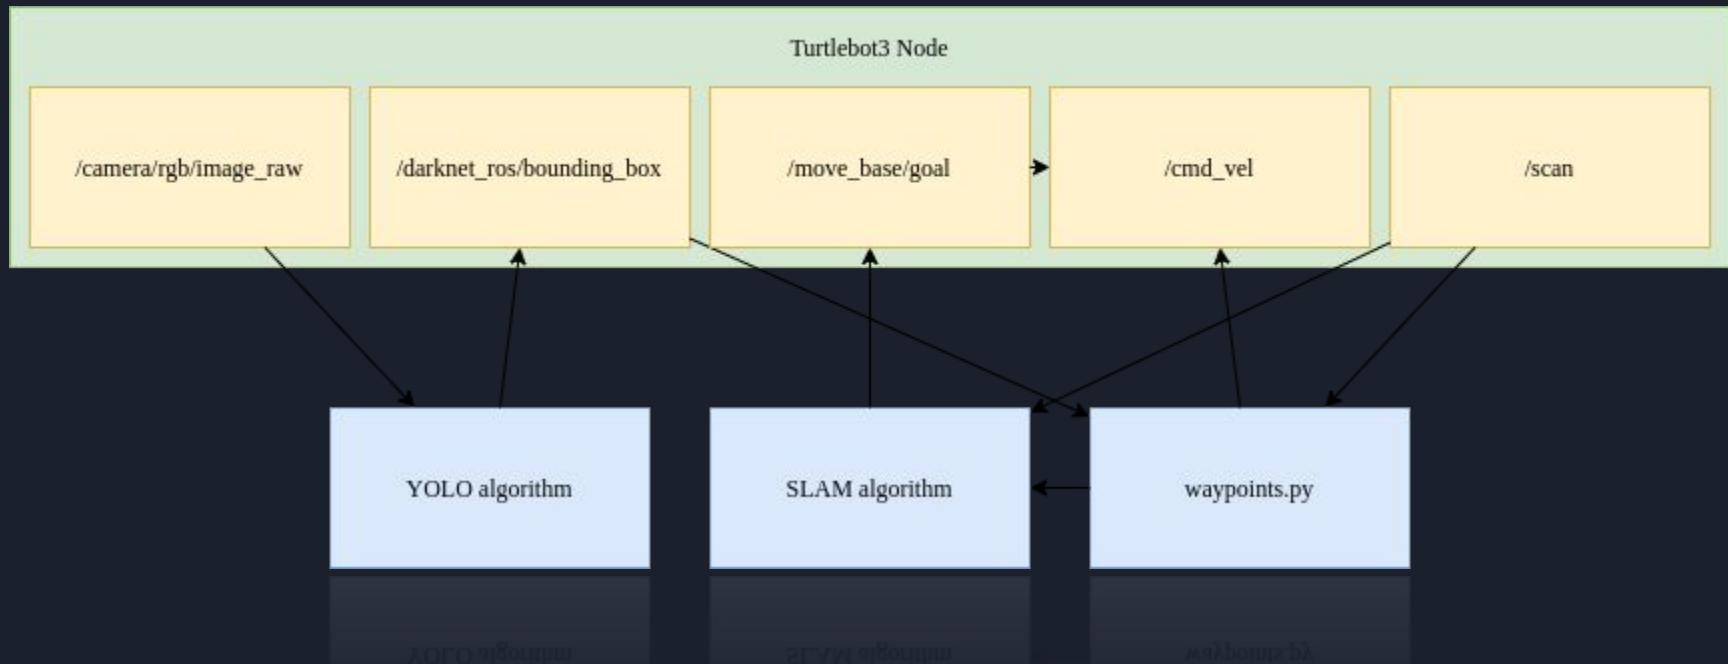


| | | | | | |
|--------------|---------------|---------------------|--------------|---------------|---------------|
| Area of Work | Patrick Bauer | Nathanael Cassagnol | Mark Pedroso | Pablo Trivino | Noah Avizemer |
|--------------|---------------|---------------------|--------------|---------------|---------------|



Current Speaker:
Patrick Bauer

ROS Communications



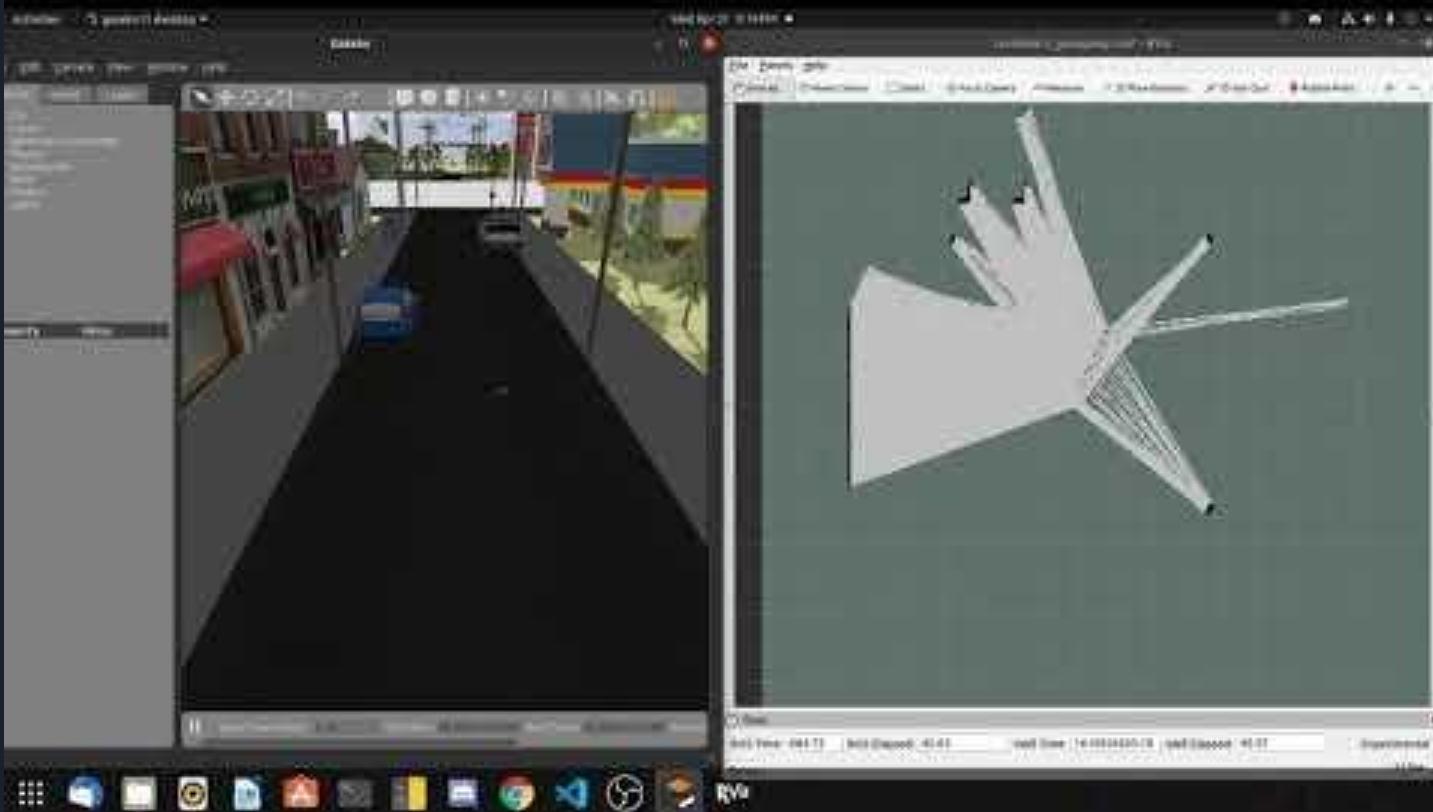


Autonomous Mapping Demo

- In order to navigate around the world, it must first be mapped out
 - Uses Gmapping SLAM
- Done autonomously with a turtlebot
 - Navigates to closest unexplored areas.
- Increased sensor range to map outdoor environment.



Current Speaker:
Noah Avizemer



Current Speaker:
Noah Avizemer

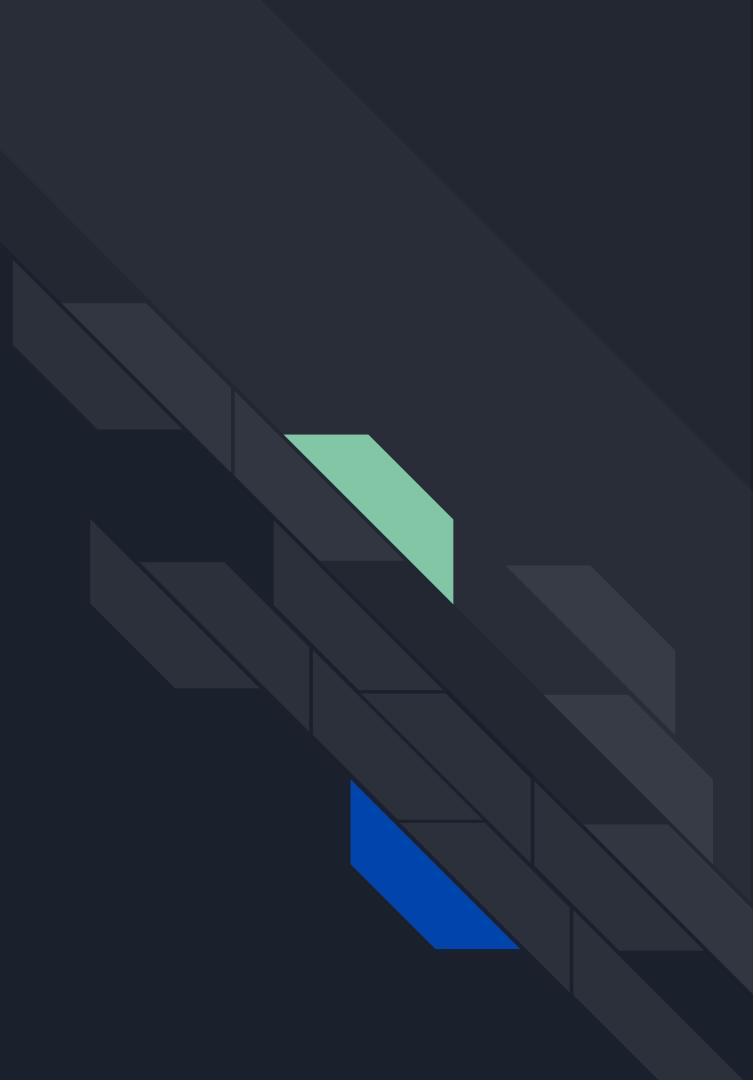


COCO Dataset Demo

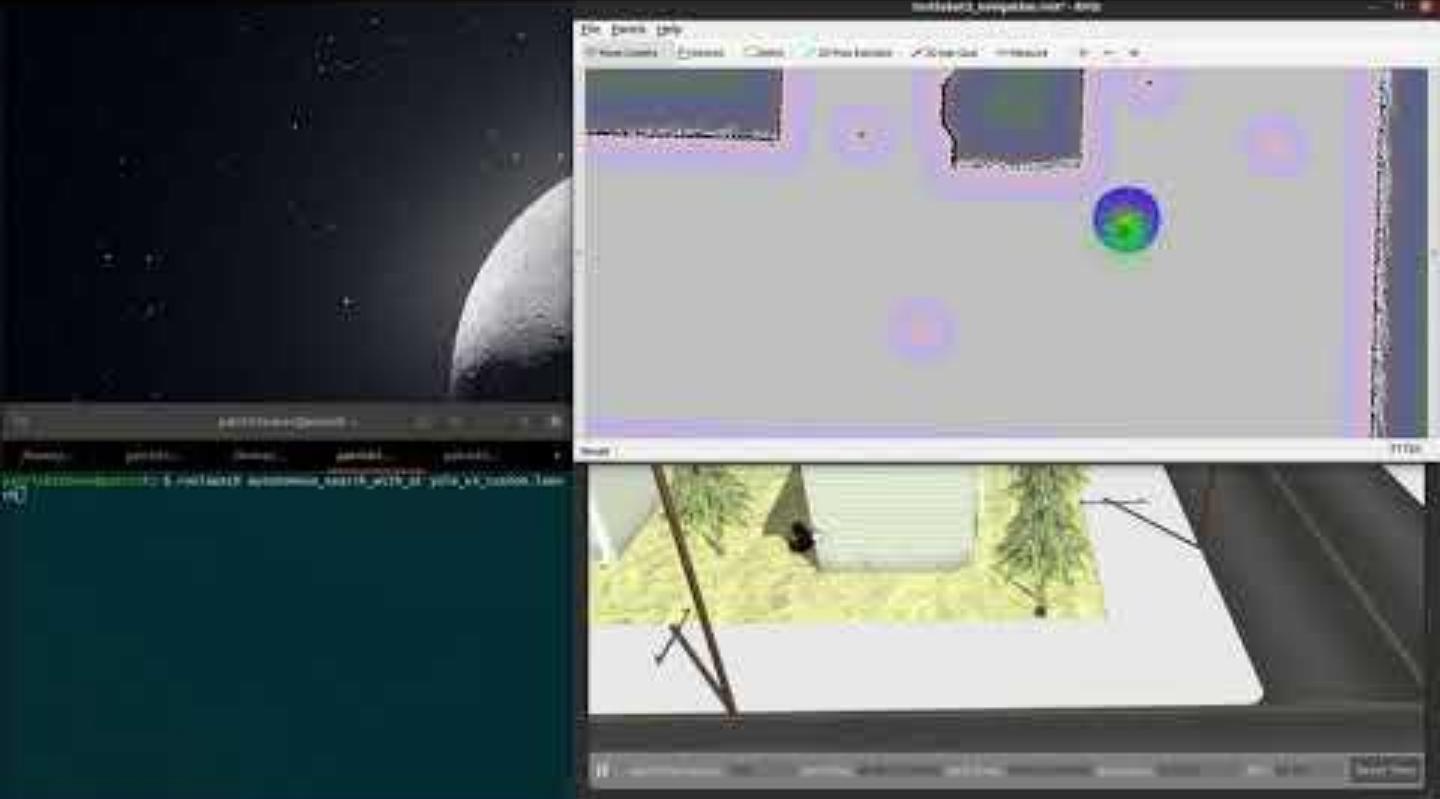


Current Speaker:
Pablo Trivino

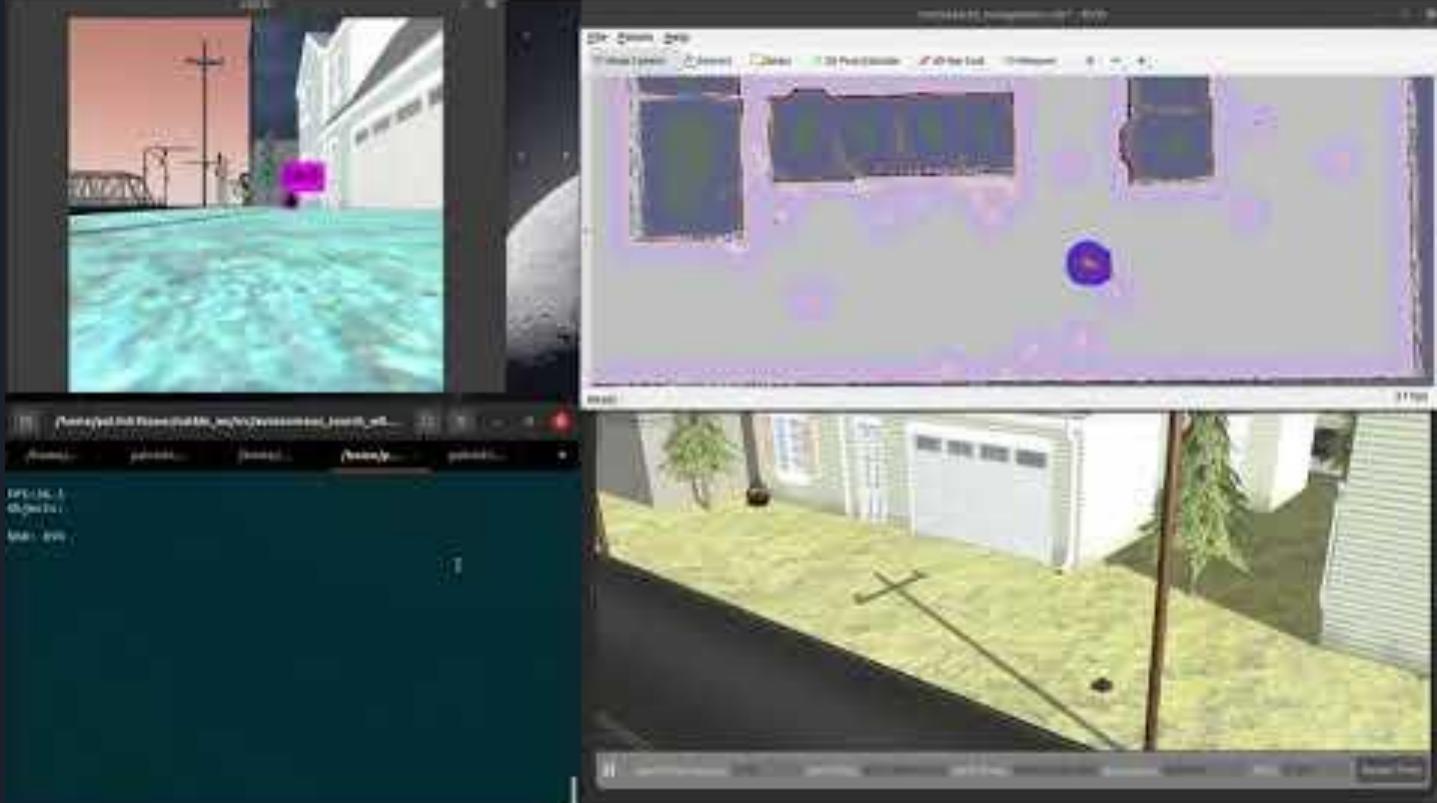
Project Demo



Current Speaker:
Patrick Bauer



Current Speaker:
Patrick Bauer



Current Speaker:
Patrick Bauer

Administration - Budget

- Took advantage of open-source software and tools
 - Free
- Purchased access to The Construct's ROS development studio
 - \$50 a month for premium version (x3)
 - \$5 a month for increased storage (x2)

TOTAL: \$160



Current Speaker:
Noah Avizemer

Administration - Teamwork

- Bi-weekly meetings with our sponsor through Skype
- Team meetings & communication through Discord
- Project management through JIRA

The screenshot shows a Jira Software board for the project 'Drone Swarm AI'. The board is titled 'DS Sprint 2' and has the following columns:

- TO DO:**
 - Find out how to implement a pathfinding algorithm within Gazebo/ROS
 - Pathfinding Algorithm
- IN PROGRESS:**
 - Create Drone & Command & Control Nodes in ROS
 - Drone Communication
 - Create Topics within ROS for node communication
 - Drone Communication
 - Create a datastructure for communicating data between parts of the drone
 - Drone Communication
 - Create a datastructure for communicating data between parts of the Command & Control Laptop
 - Drone Communication
 - Create a datastructure for communicating data between drones & the Command & Control Laptop
 - Drone Communication
 - Find a way to simulate a microwave burst within Gazebo/ROS
 - Special Illumination
- DONE:**
 - Create Gazebo Environment
 - Create Simulated Environment within Gazebo/ROS
 - Add terrain to environment
 - Create Simulated Environment within Gazebo/ROS
 - Add Streets to environment
 - Create Simulated Environment within Gazebo/ROS
 - Add Houses to environment
 - Create Simulated Environment within Gazebo/ROS
 - Add trees, cars, bushes to environment
 - Create Simulated Environment within Gazebo/ROS
 - Add drone to environment
 - Create Simulated Environment within Gazebo/ROS



Administration - Challenges

Simulation Setup:

- Dealing with different versions of ROS
- Packages not working
- Deprecated packages
- A lack of documentation on packages and methodologies
- Trouble getting started due to a lack of experience and direction

Object Detection:

- Changed target model multiple times, requiring us to re-gather training data
- Turtlebot cameras do not have 360 degree vision
- Implementing YOLOv3 on a local system
- Low FPS for object detection without GPU

Pathfinding:

- Changed requirements requiring a new approach
- The version of ROS used prevented multi-robot implementation

Administration - Deliverables

https://github.com/patricklbauer/autonomous_search_with_ai

Repository contains:

- Autonomous Search With AI package
- Readme with instructions to set up the required software and workspace.
- Final design document including our research and how the project works.

The screenshot shows the GitHub README.md page for the 'Autonomous Search With AI' repository. The page includes sections for 'About The Project', 'Built With', 'Getting Started', and 'Installation'. The 'About The Project' section mentions it's a UCF Senior Design Project sponsored by Lockheed Martin. The 'Built With' section lists ROS, Gazebo, YOLOv3, SLAM, and LabelImg. The 'Getting Started' section provides two steps: creating an Ubuntu 20.04 LTS VM and installing ROS-Noetic & Dependencies. The GitHub interface shows file navigation on the left and commit history at the bottom.



Current Speaker:
Patrick Bauer

Questions?

bb8

